Unsupervised learning (and a hint of GenAI)

Dimensional reduction, clustering, “orthogonalization”, and generation

Looking for simpler representations of our data

1. Dimensionality reduction (“compress”)
   a. Your data has N dimensions; find a compressed representation in K (K << N) dimensions
   b. Finding a “manifold” that the data lies near to
2. Clustering (“sparsify”)
3. Independent representations (“separate”)

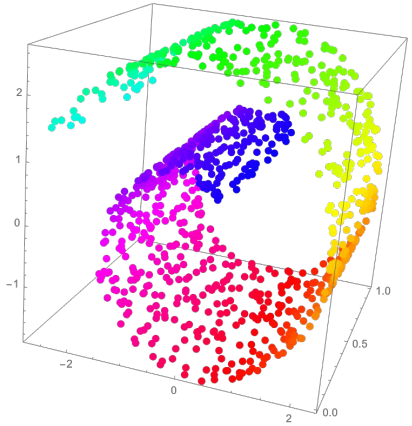

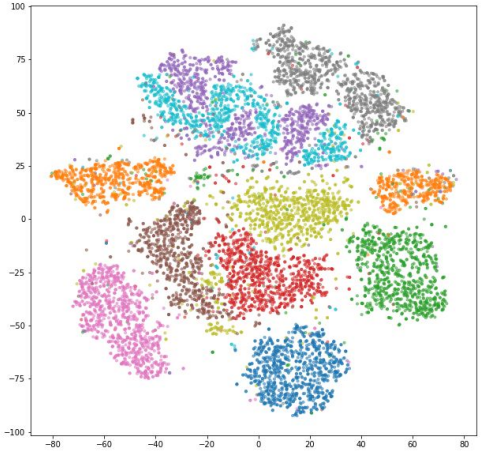

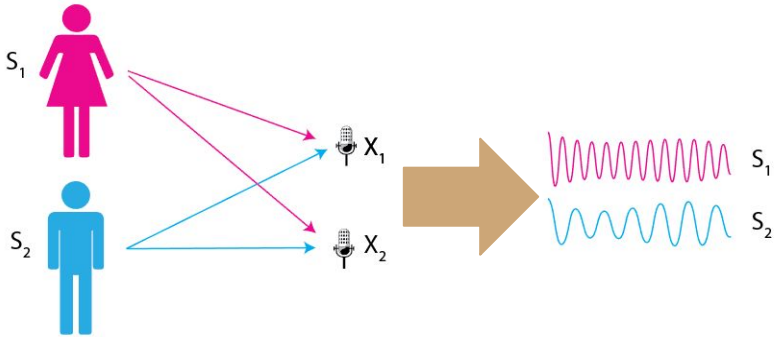

Principal component analysis (PCA)

- IDEA 💡: rotate your axes so that, in the new reference system, your data become uncorrelated
  - Example: height and weight → size and “fatness”
  - How? (Whispering): find the eigenvalues and eigenvectors of the covariance matrix
- This is (simple) *separation*
- *Compression*: keep only the k axes with the largest variance

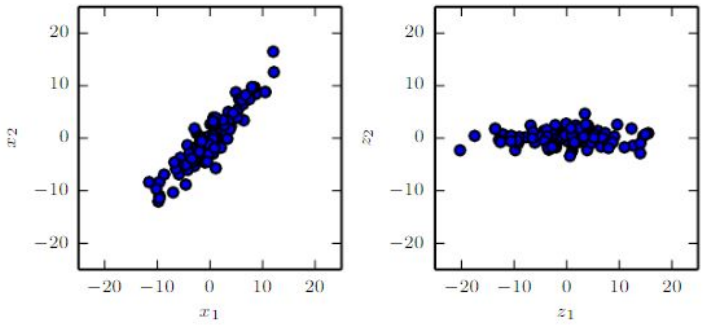

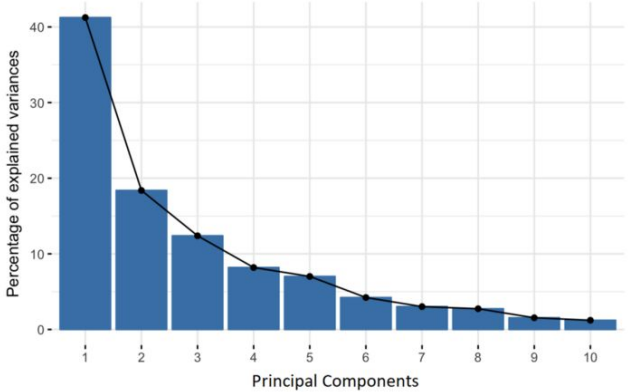

K-means clustering

https://t.ly/l9CF

- IDEA 💡: choose k random “centroids”, make them “gather” the nearest data points, re-compute the centroids, repeat till convergence
- Different initial conditions → different clusters (try several times)
- What’s the best clustering? (“average dispersion”)
- How to choose K (*a posteriori*): elbow point (also for PCA!)

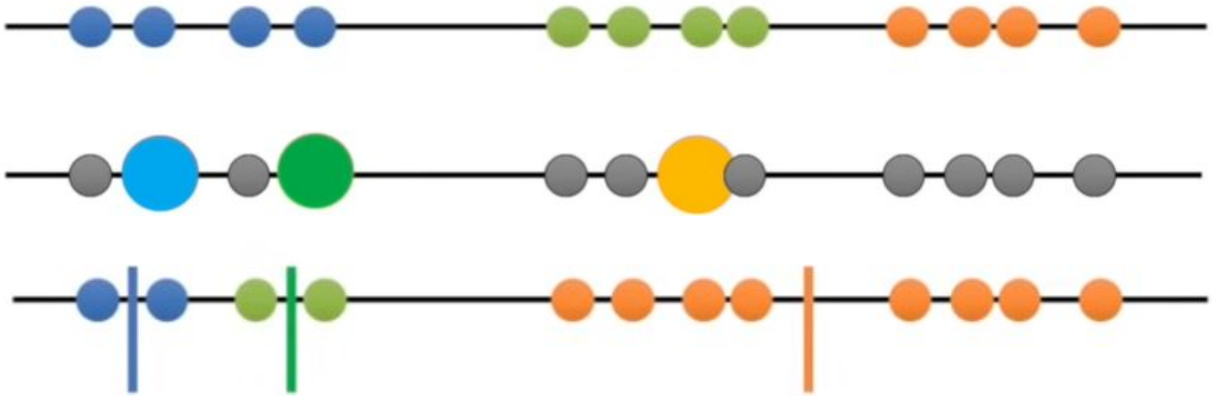

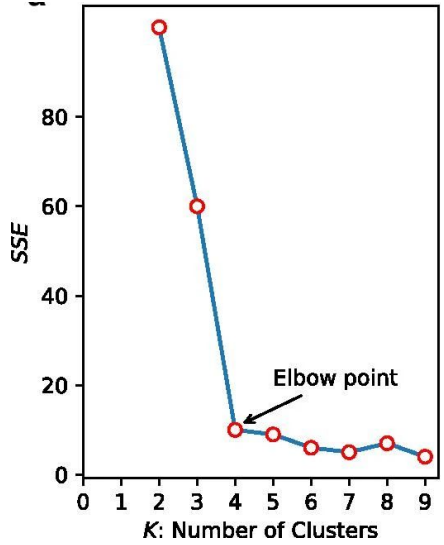

Autoencoders

- An autoencoder is a neural network that is trained to attempt to copy its input to its output...
- ... passing through a tight “information bottleneck”
- Three parts: encoder, “latent space” (or code), decoder
- “Non-linear” PCA

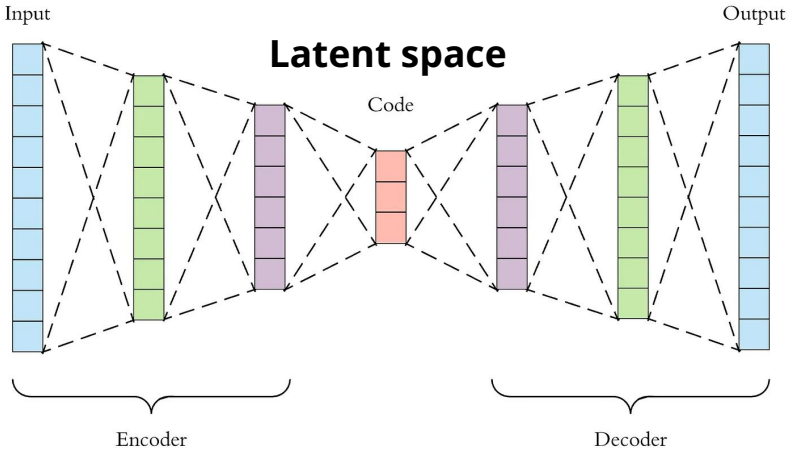

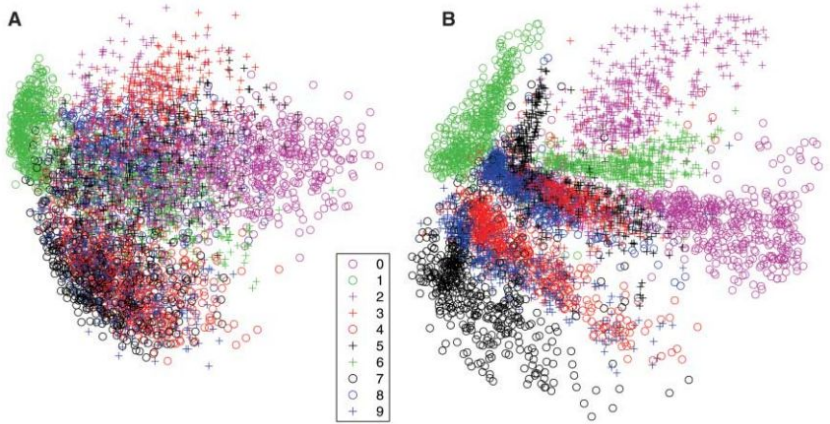

Decoding or generating?

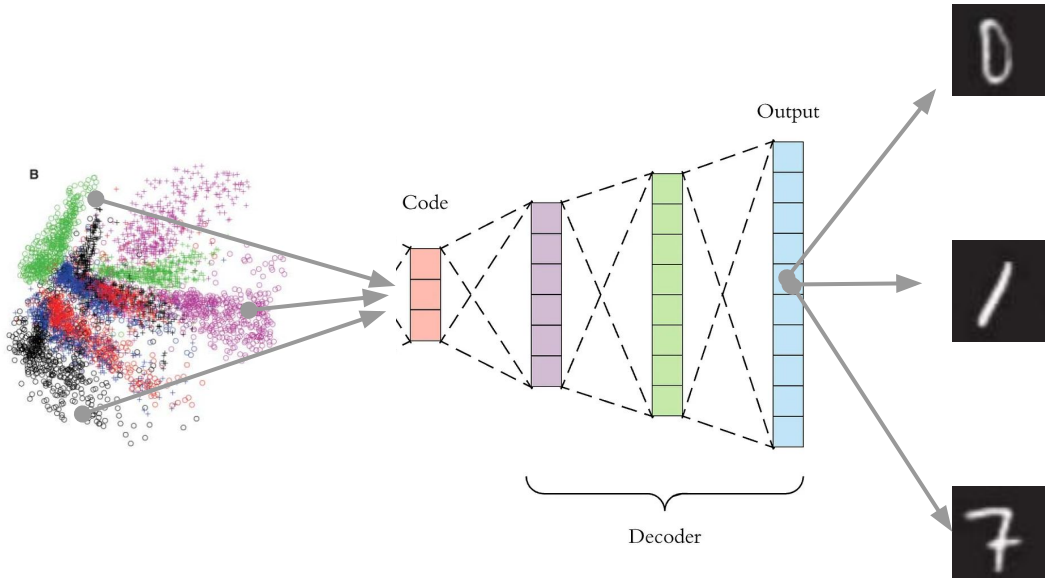

The Two Philosophies of Generative AI

Approach 1: The "Latent Blueprint"

Core Idea: Compress complex data into a rich "blueprint" (a latent space), then generate new things by creating new, original blueprints.

Approach 2: The "Predictive Chain"

Core Idea: Predict the next element in a sequence. This simple task, when mastered, leads to powerful emergent abilities.

Approach 1: "Latent Blueprint"

Starts with an Autoencoder: Compress complex data (e.g., an image) into a small, dense "latent representation" (the blueprint)

The Generative Leap: Instead of just reconstructing, **we learn to sample** new points from this latent space

Analogy: Think of the latent space as a "control panel." Each "dial" controls a feature (age, smile). Generating is like setting the dials

Models in this class: VAEs, GANs, Diffusion Models

The Magic: Shared "Blueprints"

What if we compress two modalities (like images and text) into the same shared latent space?

The model learns to map the text "a blue car" and an image of a blue car to the exact same "blueprint"

This is the key to modern text-to-image generation (DALL-E, Midjourney, Stable Diffusion)

Your Prompt (Text) → Text Encoder → Latent Blueprint → Image Decoder → Your New Image

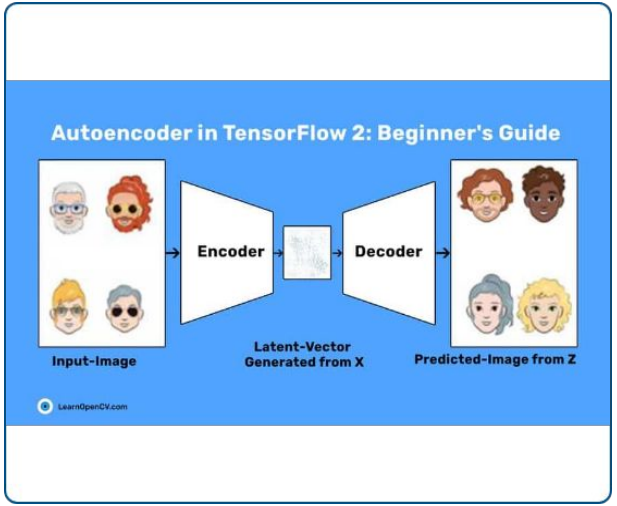

Approach 2: "Predictive Chain"

A different philosophy: The model is only trained to predict the **next element in a sequence**

This is "autocomplete on steroids"

The "Emergent" Power: To get really good at this, the model must incidentally learn grammar, context, facts, and even reasoning.

Generation: **The model predicts one word, adds it to the prompt, and repeats**, building its answer token-by-token.

This is the engine behind all Large Language Models (LLMs) like GPT.