



# xAI

eXplainable Artificial Intelligence

# What xAI is (“opening the black box”)

- Deep Neural Networks functions as “black boxes”
  - Intrinsically non-interpretable, “raw”, features (pixels, voxels, wavelets, etc.)
  - Highly non-linear (“hierarchical representations”)
- *xAI is a broad field of research in AI concerning development of methods to increase trust and understanding of ML model's predictions*
  - xAI tries to make the “box” more transparent and understandable
  - “Why did you think that's a cat?”
- The “shades-of-grey-box” problem
  - Trade-off between interpretability and representation power of ML models
  - Simple models (e.g., linear models, decision trees) interpretable, yet admittedly not very powerful
  - Complex models (like deep neural networks) powerful, yet difficult “to understand”

# What is an explanation?

- It depends on whom you ask
  - Developers ("is it working?"), expert users (e.g., doctors; "which elements influenced the decision?"), end users (e.g., patients; "what does it mean?")
  - xAI mainly skewed towards developers
- What's a good explanation?
  - (Most likely) not a simple rule
  - Visual, conceptual (text), exemplars?
  - Local vs global
  - Different methods can (and do) disagree
  - There are no clear measures to quantify the goodness of an explanation
  - (Is segmentation an explanation?)
- Is everything explainable?
  - Decision boundaries in 4+ dimensions
  - Blaise Pascal and the *esprit de finesse*
  - Split-brain patients and flashing words

# Why xAI is important

- Helps ML developers understand system outputs simply and quickly
- Can suggest solutions and spot anomalies to investigate
- Makes possible to understand why a mistake was made
  - Or was not made (correct answer for the “wrong” reason)
  - Dog in the snow = wolf
- And train the system to stop it from happening again
- Helps to meet the “right to explanation” required by the GDPR
- Helps in avoiding discrimination and biases in decisions

# Evaluating xAI

- The “human side”
  - Understandability: Can humans easily grasp the explanations provided by the xAI method?
  - Actionability: Do the explanations help users make informed decisions or take appropriate actions?
  - Trustworthiness: Do users trust the explanations and the underlying AI model?
- The “machine side”
  - Fidelity: How accurately do the explanations reflect the true workings of the AI model?
    - “Decision gradient” aligns with explanation? Counterfactuals do change the decision?
  - Sensitivity: Are explanations stable and consistent across similar inputs?
    - Perturbation analysis
  - Completeness: Do explanations cover all relevant factors?
    - Comparing explanations with domain knowledge or alternative xAI methods
- The “interface side”
  - Do humans change their attitude toward/their way of using the AI system?
  - Is it this change appropriate/useful/sound?
- It's all quite difficult to quantify!

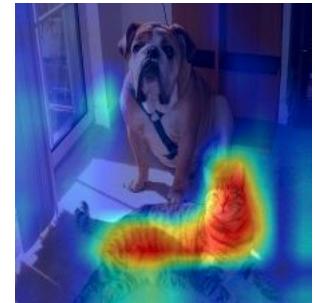
# Taxonomy of xAI methods

- Importance methods
  - What part of the input data influenced the decision (and how much)?
  - Region in an image, features, “representations”
- Exemplar methods
  - What training points impacted most the decision? - “Proponents” and “opponents”
  - Counterfactual explanations: “minimal” input change to make the machine change its mind
- Distillation methods
  - Train an interpretable (“white box”) model to reproduce the output of a complex model
  - “Structural” distillation (e.g., NN into a random forest), “complexity” distillation (large NN into smaller NN)
- Intrinsic methods (“Interpretable AI”)
  - Models that provide, alongside the main output, an explanation
  - Optimizing both model performance and a certain quality of the explanations produced
  - Prototypes, templates, concepts

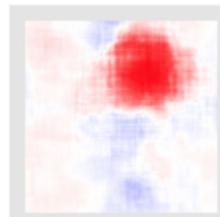
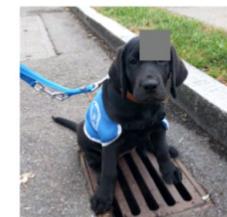
# Importance examples

- CAM (Class Activation Maps), Grad-CAM, Integrated Gradients, Smooth Gradients
  - Aim to identify which regions in the image were relevant to a specific class
  - By looking at the magnitude of the gradients (class  $C_k$ , representation  $R_i$ ) flowing through the network layers
  - Useful to measure how much each pixel/region activate the class predicted by the network
- Occlusion sensitivity
  - “Perturbing” (masking) the input
  - And looking at the effects on classification
  - Saliency = output variation

$$\frac{\partial C_k}{\partial R_i} \rightarrow$$



predicted class: cat



class dog

# Exemplar example: Gradient tracing/TrackIN

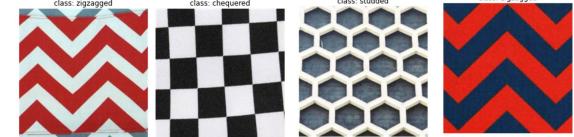
- Data influence: How much has this training point influenced the prediction for this test point?
- Retrain the model without the training point and test
- This is prohibitively expensive!
- Resorting to approximations

$$\sum_e \nabla_{\mathbf{w}_e} \text{loss}(x_{\text{train}}) \cdot \nabla_{\mathbf{w}_e} \text{loss}(x_{\text{test}})$$

Test point



Training points



Test image



church



church

Proponents



church

Opponents



castle



castle



castle



af-chameleon



af-chameleon



af-chameleon



brocoli



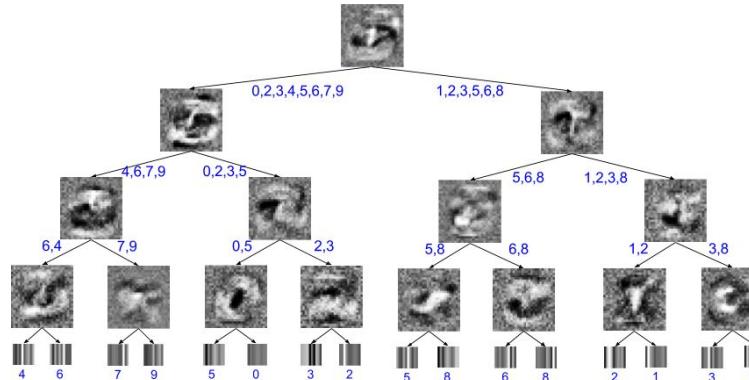
agama



jackfruit

# Distillation example

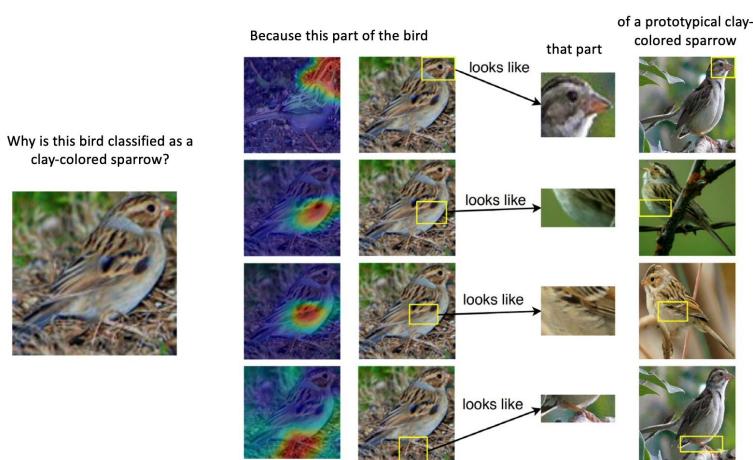
- DNN into a “soft decision tree”
- Each node is a “learned representation”
- Explanation = most probable path
- Why distillation?
  - Simply, the DNN is more powerful
  - And you can build better decision trees by using “**soft label**” generated by the DNN



# Intrinsic example

- Heat-maps look good, don't they?
  - Explanation seems more like "attention", not related to the specific class
  - They can be quite "unsensitive" (distillation can help with that)
- "This looks like that" — prototypes
  - The network dissects the image by finding prototypical parts
  - And combines evidence from the prototypes to make a final classification
  - Explanation mimicking (in part) the way an ornithologist would reason

	Test Image	Evidence for Animal Being a Siberian Husky	Evidence for Animal Being a Transverse Flute
Explanations Using Attention Maps			



# Extreme gradient boosting (XGBoost)



# Why XGBoost?

*The XG Boosting algorithm uses advanced regularization techniques to suppress weights, prevent overfitting, and enhance its performance in real-world scenarios. On top of this, the implementation allows the algorithm to cache data and utilize multiple CPU cores for speedy processing. The enhanced performance and speed have made **XGBoost one of the most popular machine-learning algorithms in recent years.***

# What is XGBoost?

XGBoost gives a prediction function  $f$ , by regressing a set of parameters  $\theta$  on a dataset of predictors  $\mathbf{X}$  related to the dependent variable of interest  $y$ .

Function  $f$  is non-linear and admits general interactions between predictors.

$$y = f(\mathbf{X} | \theta)$$

# Ensemble ML

- “Weak” learners: simple models (simple = does not overfit)
- Ensembles of weak learners can be very strong
- Bagging (parallel)
  - Each weak learner “looks” only at part of the features
  - We average (or count the “votes”) the answers of the weak learners
  - Each weak learner is trained independently
- Boosting (sequential)
  - The first weak learner approximates  $\hat{y}^{(1)} = f^{(1)}(x) \approx y$
  - The second one learns  $\hat{y}^{(2)} = y - \hat{y}^{(1)}$
  - The  $t$ -th learner learns  $\hat{y}^{(t)} = y - \hat{y}^{(t-1)}$
  - “Gradient boosting” is a generalisation

$$\hat{y}_i^{(0)} = 0$$

$$\hat{y}_i^{(1)} = f_1(x_i) = \hat{y}_i^{(0)} + f_1(x_i)$$

$$\hat{y}_i^{(2)} = f_1(x_i) + f_2(x_i) = \hat{y}_i^{(1)} + f_2(x_i)$$

...

$$\hat{y}_i^{(t)} = \sum_{k=1}^t f_k(x_i) = \hat{y}_i^{(t-1)} + f_t(x_i)$$

$$\begin{aligned}\text{obj}^{(t)} &= \sum_{i=1}^n (y_i - (\hat{y}_i^{(t-1)} + f_t(x_i)))^2 + \sum_{i=1}^t \omega(f_i) \\ &= \sum_{i=1}^n [2(\hat{y}_i^{(t-1)} - y_i)f_t(x_i) + f_t(x_i)^2] + \omega(f_t) + c \\ &= 2 \sum_{i=1}^n [g_i j_t(x_i) + \frac{1}{2} h_i j_t^2(x_i)] + \omega(f_t) + c\end{aligned}$$

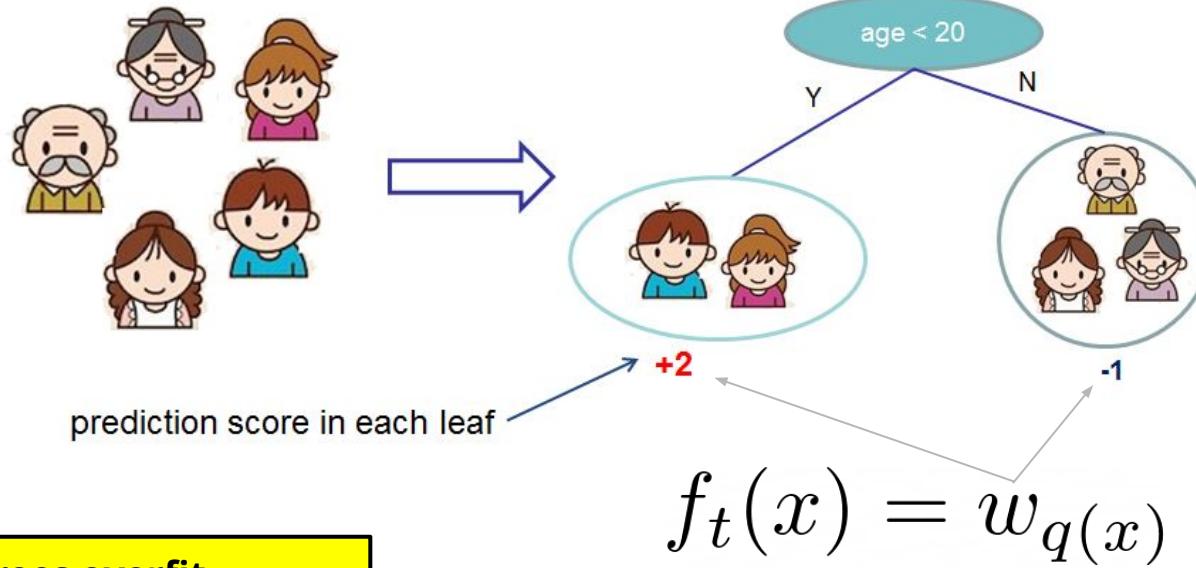
**g** stands for gradient

**h** stands for... second derivative (Hessian)

# Trees

Input: age, gender, occupation, ...

Like the computer game X



Deep trees overfit  
Shallow trees are weak learners

# Learning the leaf weights

$$\omega(f) = \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2$$

Number of leafs

Regularisation

Best solution

$$\begin{aligned}\text{obj}^{(t)} &= 2 \sum_{i=1}^n [g_i w_{q(x_i)} + \frac{1}{2} h_i w_{q(x_i)}^2] + \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2 \\ &= 2 \sum_{j=1}^T \left[ \left( \sum_{i \in I_j} g_i \right) w_j + \frac{1}{2} \left( \sum_{i \in I_j} h_i + \lambda \right) w_j^2 \right] + \gamma T \\ &= 2 \sum_{j=1}^T [G_j w_j + \frac{1}{2} (H_j + \lambda) w_j^2] + \gamma T\end{aligned}$$

$$\begin{aligned}w_j^* &= -\frac{G_j}{H_j + \lambda} \\ \text{obj}^* &= -\frac{1}{2} \sum_{j=1}^T \frac{G_j^2}{H_j + \lambda} + \gamma T\end{aligned}$$

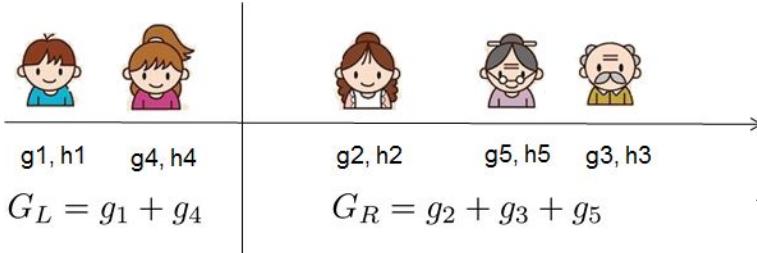
# Learning the tree structure

This formula can be decomposed as:

1. The score on the new left leaf
2. The score on the new right leaf
3. The score on the original leaf
4. Regularization on the additional leaf

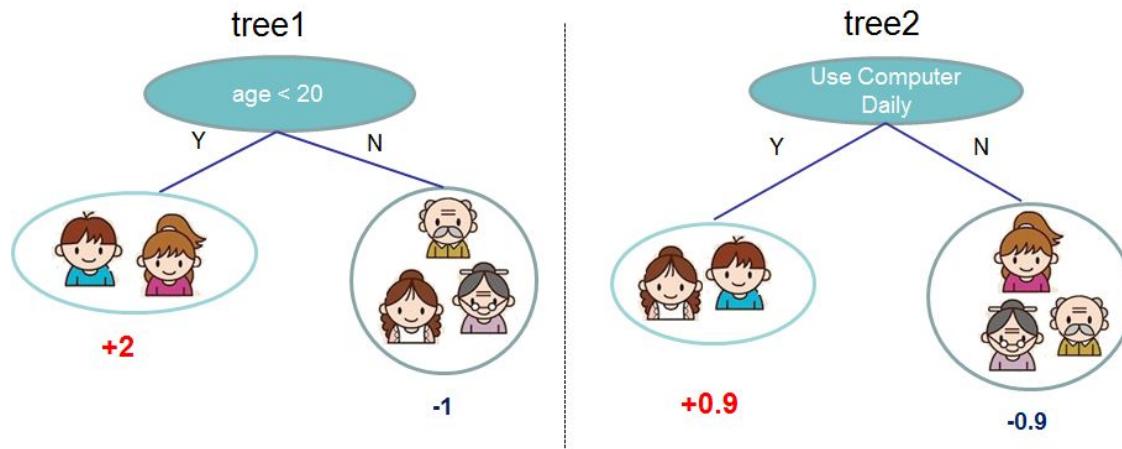
$$Gain = \frac{1}{2} \left[ \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{(G_L + G_R)^2}{H_L + H_R + \lambda} \right] - \gamma$$

If the gain is smaller than  $\gamma$ , we would do better not to add the new branch (this is usually called "pruning")



A left to right scan is sufficient to find the best split

# Forests



$$f(\text{boy}) = 2 + 0.9 = 2.9$$

$$f(\text{old man}) = -1 - 0.9 = -1.9$$

**Bagging:**  
**Random Forest**

**Boosting:**  
**Boosted trees**

# What about “extreme” in XGBoost?

*“XGBoost is an **optimized distributed** gradient boosting library designed to be **highly efficient, flexible, and portable**. It implements machine learning algorithms under the Gradient Boosting framework. XGBoost provides a parallel tree boosting that solve many data science problems in a **fast and accurate** way. The same code runs on major distributed environment and **can solve problems beyond billions of examples**”*

**There are other gradient boosting libraries,  
such as LightGBM and CatBoost**



SHAP



# What is SHAP?

Here SHAP is just an algorithm that, given the function  $f$  returned by XGBoost and given one individual  $i$  (one set of values for the predictors  $\mathbf{X}_i$  - for example: age 37, BMI 28.4, systolic blood pressure 135), estimates the contribution of each predictor to the specific prediction  $y_i$ .

In simple words, SHAP gives a principled estimation of, for example, **how much the prediction  $y_i$  would change if the age of individual  $i$  were not observed**

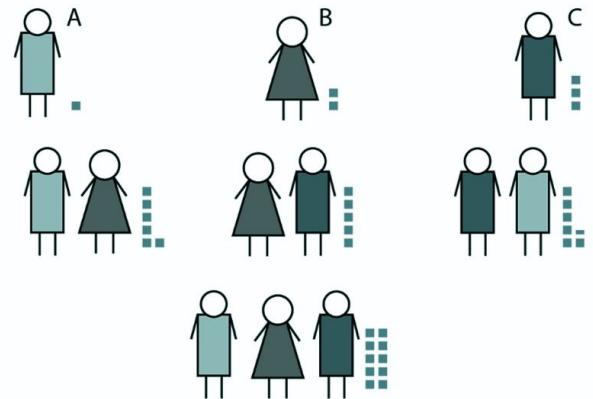
$$\Delta y_i = \text{SHAP}(\text{age}_i)$$

$$\text{SHAP}(\text{age}_i) = f(\mathbf{X}_i | \theta) - f(\mathbf{X}_i \setminus \text{age}_i | \theta)$$

**SHAP is an Importance Method**

# Shapley value

- A game theory concept introduced by Lloyd Stowell Shapley (Nobel Prize in economic sciences, 2012)
- **The Shapley value gives a “fair” estimate of the contribution of each player participating in a collaborative work**
- There is a game, there are  $N$  players; for each possible team  $T$  of  $k$  ( $\leq N$ ) players there is a payoff  $V(T)$  (that measures the “value of the team”)
- Knowing  $V(T)$  for all the possible teams, how can we measure the “value of each single player”?



$$\begin{aligned}\varphi_i(v) &= \sum_{S \subseteq N \setminus \{i\}} \frac{|S|! (n - |S| - 1)!}{n!} (v(S \cup \{i\}) - v(S)) \\ &= \sum_{S \subseteq N \setminus \{i\}} \binom{n}{1, |S|, n - |S| - 1}^{-1} (v(S \cup \{i\}) - v(S))\end{aligned}$$

# SHAP

- In SHAP the players are the predictors
- You have a predictive function/model (XGBoost, in our case)  $f(\mathbf{X} | \theta)$  and a data point  $\mathbf{X}$ , you want to analyse
- The value  $V(T)$  of a team  $T$  (a subset of the predictors) is the prediction given by  $f$  when you (conditionally) average over the predictors NOT in  $T$
- Very hard to compute exactly. SHAP makes some simplifying assumptions
- These simplifications make SHAP very efficient to compute for Decision Trees (the atomic element of XGBoost) → **TreeSHAP**
- There exist many instantiations of SHAP, for different machine learning algorithms

# Individualists, synergies, and conflicts

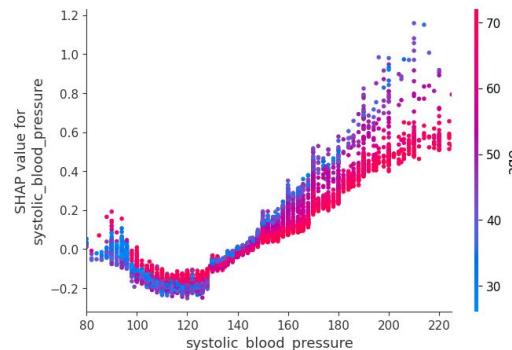
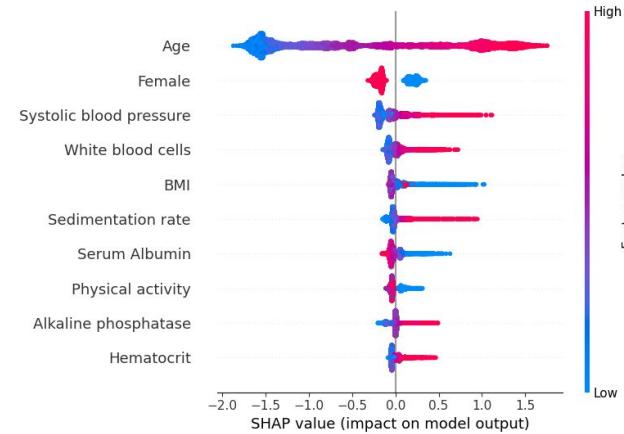
- An individualistic player C is one that gives the same “boost”  $v_C$  to every team ( $v_{AC} = \mathbf{0}$  for all players A, with  $A \neq C$ )
- Synergy: player A adds  $v_A$ ; player B adds  $v_B$ ;  $V(A + B) > v_A + v_B$ 
  - A and B boost each other game ( $v_{AB} > \mathbf{0}$ )
- Conflict: player A adds  $v_A$ ; player B adds  $v_B$ ;  $V(A + B) < v_A + v_B$ 
  - A and B compete for the same “role” ( $v_{AB} < \mathbf{0}$ )
- **Main effect** and **interaction values**

$$V(T) = \sum_{A \in T} v_A + \frac{1}{2} \sum_{(A, B) \in T} v_{AB}$$

$$\text{SHAP}_A = v_A + \frac{1}{2} \sum_B v_{AB}$$

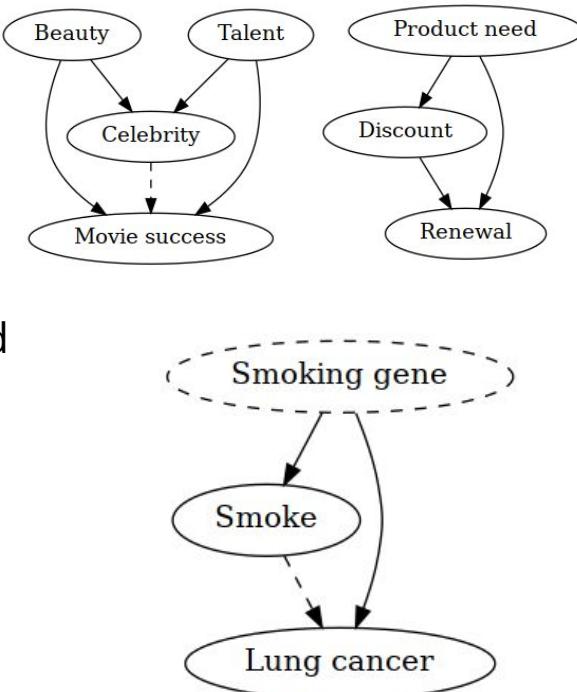
# A quick example on mortality

- Age is by far the most significant predictor
- Sex (second place) acts largely as a constant factor
- There is, quite consistently, a monotonic relationship between feature value (e.g., Age) and the effect on prediction (SHAP)
- There is a strong “interaction” between age and SBP



# Explanation is not causation

- AI doesn't try to build a "physical model" of the data
- AI just "wants to predict"
- AI can discover, and take advantage of, highly complex and non-linear "correlations" among predictors, and between predictors and response
  - A good collider can be better than the real causes (and more parsimonious)
  - Direct and mediated effects entangled in explanations
- xAI just unearths these "correlations on steroids"
  - Don't forget that you are explaining your model, not reality
- And, of course, unobserved confounders are the hardest problem with or without AI



# Yet...

- It is difficult to dismiss the predictive power of AI
  - Prediction is important *per se* (Where and when can I intervene most effectively?)
  - And there are anyhow many factors that we cannot control, but are determinant
- And the patterns of “influence” uncovered by SHAP are just too intriguing
- At the very least, AI + xAI represents today one of the most promising instruments for scientific *exploration*
- As for *explanation*:
  - Strongly nonlinear causal effects (e.g., U-shaped, one predictor “gating” the effects of another) will be mostly under-detected and under-measured by traditional methods
  - xAI can make the most out your prior knowledge
  - There are attempts to make AI (and xAI) more “causality-aware”
  - **Double Machine Learning** tries to extract causal knowledge from observational data
- But we leave that for a second course...

Ci vediamo giovedì 27 novembre  
(parliamo di esame?)